

# CUSTOMER CHURN PREDICTION

In [106]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score, recall_score, f1_score
from sklearn import svm
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

In [2]:

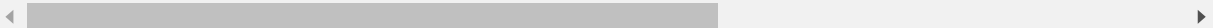
```
data = pd.read_csv('C:\Churn_Modelling.csv')
```

In [3]:

```
data.head()
```

Out[3]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	N
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	
3	4	15701354	Boni	699	France	Female	39	1	0.00	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	

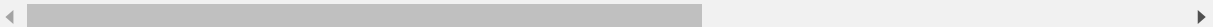


In [4]:

```
data.tail()
```

Out[4]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	N
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	



In [5]:

```
data.shape
```

Out[5]: (10000, 14)

In [6]:

```
print("Number of Rows", data.shape[0])
print("Number of Columns", data.shape[1])
```

Number of Rows 10000  
Number of Columns 14

In [7]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   RowNumber              10000 non-null  int64  
1   CustomerId             10000 non-null  int64  
2   Surname                10000 non-null  object  
3   CreditScore            10000 non-null  int64  
4   Geography              10000 non-null  object  
5   Gender                 10000 non-null  object  
6   Age                    10000 non-null  int64  
7   Tenure                 10000 non-null  int64  
8   Balance                10000 non-null  float64 
9   NumOfProducts         10000 non-null  int64  
10  HasCrCard              10000 non-null  int64  
11  IsActiveMember        10000 non-null  int64  
12  EstimatedSalary        10000 non-null  float64 
13  Exited                 10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [10]:

```
data.isnull().sum()
```

```
Out[10]: RowNumber      0
CustomerId    0
Surname       0
CreditScore   0
Geography     0
Gender        0
Age           0
Tenure        0
Balance       0
NumOfProducts 0
HasCrCard     0
IsActiveMember 0
EstimatedSalary 0
Exited        0
dtype: int64
```

In [12]:

```
data.describe(include = 'all')
```

```
Out[12]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	
<b>count</b>	10000.00000	1.000000e+04	10000	10000.000000	10000	10000	10000.000000	10000
<b>unique</b>	NaN	NaN	2932	NaN	3	2	NaN	
<b>top</b>	NaN	NaN	Smith	NaN	France	Male	NaN	
<b>freq</b>	NaN	NaN	32	NaN	5014	5457	NaN	
<b>mean</b>	5000.50000	1.569094e+07	NaN	650.528800	NaN	NaN	38.921800	
<b>std</b>	2886.89568	7.193619e+04	NaN	96.653299	NaN	NaN	10.487806	
<b>min</b>	1.00000	1.556570e+07	NaN	350.000000	NaN	NaN	18.000000	
<b>25%</b>	2500.75000	1.562853e+07	NaN	584.000000	NaN	NaN	32.000000	
<b>50%</b>	5000.50000	1.569074e+07	NaN	652.000000	NaN	NaN	37.000000	
<b>75%</b>	7500.25000	1.575323e+07	NaN	718.000000	NaN	NaN	44.000000	

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	
max	10000.00000	1.581569e+07	NaN	850.000000	NaN	NaN	92.000000	1

In [13]: `data.columns`

Out[13]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited'], dtype='object')

In [15]: `data = data.drop(['RowNumber', 'CustomerId', 'Surname'],axis = 1)`

In [16]: `data.head()`

Out[16]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

In [17]: `data['Geography'].unique()`

Out[17]: array(['France', 'Spain', 'Germany'], dtype=object)

In [19]: `data = pd.get_dummies(data,drop_first = True)`

In [20]: `data.head()`

Out[20]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	42	2	0.00	1	1	1	101348.8
1	608	41	1	83807.86	1	0	1	112542.5
2	502	42	8	159660.80	3	1	0	113931.5
3	699	39	1	0.00	2	0	0	93826.6
4	850	43	2	125510.82	1	1	1	79084.1

In [80]: `(data[data['Exited'] == 1].shape[0] / data.shape[0])*100`

Out[80]: 20.369999999999997

In [22]: `data['Exited'].value_counts()`


```
Out[22]: 0    7963
         1    2037
         Name: Exited, dtype: int64
```

```
In [31]: sns.countplot(data['Exited'])
```

C:\Users\SYS\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


```
warnings.warn(
```

```
Out[31]: <AxesSubplot:xlabel='Exited', ylabel='count'>
```



```
In [82]: plt.figure(figsize=(8,5))
         sns.kdeplot(x = data['Tenure'], hue = data['Exited'])
```

```
Out[82]: <AxesSubplot:xlabel='Tenure', ylabel='Density'>
```



```
In [40]: x = data.drop('Exited',axis = 1)
         y = data['Exited']
```

```
In [33]: y
```

```
Out[33]: 0      1
          1      0
          2      1
          3      0
          4      0
          ..
          9995    0
          9996    0
          9997    1
          9998    1
          9999    0
          Name: Exited, Length: 10000, dtype: int64
```

```
In [44]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size = 0.20,random_stat
```

```
In [47]: sc = StandardScaler()
```

```
In [48]: x_train = sc.fit_transform(x_train)
          x_test = sc.transform(x_test)
```

```
In [49]: x_train
```

```
Out[49]: array([[ 1.058568 ,  1.71508648,  0.68472287, ..., -0.57831252,
                  -0.57773517,  0.90750738],
                 [ 0.91362605, -0.65993547, -0.6962018 , ...,  1.72916886,
                  -0.57773517,  0.90750738],
                 [ 1.07927399, -0.18493108, -1.73189531, ...,  1.72916886,
                  -0.57773517, -1.10191942],
                 ...,
                 [ 0.16821031, -0.18493108,  1.3751852 , ..., -0.57831252,
                  -0.57773517, -1.10191942],
                 [ 0.37527024, -0.37493284,  1.02995403, ..., -0.57831252,
                  1.73089688,  0.90750738],
                 [ 1.56586482,  1.14508121,  0.68472287, ..., -0.57831252,
                  1.73089688,  0.90750738]])
```

## LogisticRegression

```
In [51]: log = LogisticRegression()
```

```
In [52]: log.fit(x_train, y_train)
```

```
Out[52]: LogisticRegression()
```

```
In [89]: y_pred1 = log.predict(x_test)
```

```
In [90]: accuracy_score(y_test,y_pred1)
```

```
Out[90]: 0.808
```

```
In [91]: precision_score(y_test, y_pred1)
```

Out[91]: 0.5891472868217055

```
In [92]: recall_score(y_test, y_pred1)
```

Out[92]: 0.18673218673218672

```
In [61]: f1_score(y_test, y_pred1)
```

Out[61]: 0.2835820895522388

```
In [ ]: pc = TP / (FP + TP)
```

$rc = TP / (TP + FN)$

## SVM

```
In [66]: svm = svm.SVC()
```

```
In [67]: svm.fit(x_train,y_train)
```

Out[67]: SVC()

```
In [68]: y_pred2 = svm.predict(x_test)
```

```
In [93]: accuracy_score(y_test,y_pred2)
```

Out[93]: 0.861

```
In [94]: precision_score(y_test, y_pred2)
```

Out[94]: 0.8341968911917098

## KNeighbors Classifier

```
In [73]: knn = KNeighborsClassifier()
```

```
In [95]: knn.fit(x_train,y_train)
```

Out[95]: KNeighborsClassifier()

```
In [75]: y_pred3 = knn.predict(x_test)
```

```
In [96]: accuracy_score(y_test,y_pred3)
```

Out[96]: 0.824

```
In [97]: precision_score(y_test, y_pred3)
```

Out[97]: 0.6222222222222222

## Decision Tree Classifier

```
In [84]: dt = DecisionTreeClassifier()
```

```
In [85]: dt.fit(x_train,y_train)
```

Out[85]: DecisionTreeClassifier()

```
In [86]: y_pred4 = dt.predict(x_test)
```

```
In [98]: accuracy_score(y_test,y_pred4)
```

Out[98]: 0.7865

```
In [99]: precision_score(y_test, y_pred4)
```

Out[99]: 0.4772727272727273

## RandomForestClassifier

```
In [101... rf = RandomForestClassifier()
```

```
In [102... rf.fit(x_train,y_train)
```

Out[102... RandomForestClassifier()

```
In [103... y_pred5 = dt.predict(x_test)
```

```
In [104... accuracy_score(y_test,y_pred5)
```

Out[104... 0.7865

```
In [105... precision_score(y_test, y_pred5)
```

Out[105... 0.4772727272727273

## GradientBoostingClassifier

```
In [107... gbc = GradientBoostingClassifier()
```

```
In [108... gbc.fit(x_train,y_train)
```

```
Out[108... GradientBoostingClassifier()
```

```
In [109... y_pred6 = dt.predict(x_test)
```

```
In [110... accuracy_score(y_test,y_pred6)
```

```
Out[110... 0.7865
```

```
In [111... precision_score(y_test, y_pred6)
```

```
Out[111... 0.4772727272727273
```

```
In [112... final_data = pd.DataFrame({'Models': ['LR','SVC','KNN','DT','RF','GBC'],'ACC':[accu
```

```
In [113... final_data
```

```
Out[113... 

|   | Models | ACC    |
|---|--------|--------|
| 0 | LR     | 0.8080 |
| 1 | SVC    | 0.8610 |
| 2 | KNN    | 0.8240 |
| 3 | DT     | 0.7865 |
| 4 | RF     | 0.7865 |
| 5 | GBC    | 0.7865 |


```


```
In [114... sns.barplot(final_data['Models'],final_data['ACC'])
```

C:\Users\SYS\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[114... <AxesSubplot:xlabel='Models', ylabel='ACC'>
```





```
In [116... final_data = pd.DataFrame({'Models': ['LR', 'SVC', 'KNN', 'DT', 'RF', 'GBC'], 'PRE': [preci
```

```
In [118... final_data
```

```
Out[118...
  Models  PRE
0    LR  0.589147
1    SVC  0.834197
2    KNN  0.622222
3     DT  0.477273
4     RF  0.477273
5     GBC  0.477273
```

```
In [119... sns.barplot(final_data['Models'], final_data['PRE'])
```

C:\Users\SYS\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[119... <AxesSubplot:xlabel='Models', ylabel='PRE'>
```

